

Due to a misunderstanding, I originally assumed all errors should only be thrown on write commits, thus giving me a fair amount of work fixing up my reference monitor. The first, and easiest, fix, was to throw a `RepyArgumentError` whenever `wreat` was called with a negative or non-integer offset, or with non-string data.

I added an `is_open` boolean, initialized as `True` under the `LPFile` constructor, to keep track of file state. `Wreat` and `close` check this boolean on every operation and throw a `FileClosedError` if false. `Close` also writes any valid pending data to the file before setting the `is_open` boolean to false and closing the file. `Undo` will not execute if the file is already closed. This avoids writing to or closing already closed files.

Keeping track of the file length was a bit trickier. I included a `file_length` attribute, initialized to the length of the opened file in the constructor. I also save a `prev_length` attribute, initialized to the same length in the constructor. If the sum of the offset and the length of the data is greater than or equal to the current `file_length`, `prev_length` is set to the current `file_length` and `file_length` is updated to `offset + len(data)` in `wreat`. Updating `file_length` in this fashion avoids having to call `readat`, which would slow down the reference monitor's performance. The equality condition catches a corner case where, e.g., two writes of equal length are called after opening an empty file, then `undo` is called. If the file length only updated when the offset + the length of the pending data were strictly greater, `undo` would erroneously update `file_length` back to 0. Keeping track of `prev_length` allows `undo` to function correctly when 2 writes extend the length of the file, but the second write overwrites some part of the preceding write which extended beyond the initial length. `RepyArgumentError`, `FileClosedError`, `SeekPastEndOfFileError` are all thrown in order of priority as spelled out in the `repy` docs.

I also included a write lock for safe, concurrent functionality. The write lock is used on `wreat`, `undo`, and `close`, as none of these functions should be able to run in parallel, but this allows `readat` to be called in parallel and increase the reference monitor's performance. The locks are acquired, and all following functionality is encapsulated in a `try` block with lock release guaranteed in a `finally` block. Note that the write call in `close` is a direct API call to avoid deadlock.

As a sidenote, I offer the suggestion that the next iteration of this assignment include mandatory, succinct commenting, and the possibility of throwing out testcases with unclear comments or an unnecessary level of test complexity (e.g. a dozen successive writes when 2 or 3 would catch the relevant corner case). Even as a toothless threat, this may encourage best practices. I believe this would help both the students and the grading team save a fair amount of time interpreting some of the worst code etiquette offenders.