

ISP Assignment 2.3 - Learning from mistakes

My reference monitor had 3 mistakes that were caught by 5 people's attack cases.

The issues related to *race conditions* caused by improper locking, *incorrect order of raising exceptions* in case of issues, and incorrect understanding of the *writeln assignment spec*.

1. **Race Conditions:** I had not used locking within my reference monitor program during initial development. This led to attackcases exploiting my program by running multiple threads that did read, write, undo operations on a file in a loop. During some of the iterations of those loops, the state I maintained within my LPFile object became inconsistent after it was correctly checked, and the original Repry API operations were called with the inconsistent data. This led to errors being generated.

For example, my writeln function does processing to ensure that the pending_data is correct before passing it to the original Repry writeln. But if undo is called after checking the data but before the actual call to original writeln, it destroys the data that is passed (changes it to None) and thus the original writeln throws an exception.

This is not acceptable and if the state of the variables was locked between start and end of my writeln implementation, undo would not have executed in the middle and destroyed valid data. My new implementation takes care of locking.

Two attacks did correct threading attacks to fail my reference monitor.

2. **Exception Ordering Error:** My writeln function did argument and state checking in the order of RepryArgumentError, SeekPastEndOfFileError and then FileClosedError. The actual Repry api diverges from this behavior and does checking in a different order - RepryArgumentError first, FileClosedError second, SeekPastEndOfFileError third. For example, if writeln is called on a closed file with offset past EOF, Repry API states that FileClosed error should be raised but my implementation raised SeekPastEndOfFileError error which is wrong. This difference in implementation led to unexpected exceptions being generated that were not caught by the testcase handlers and bubbled up to generate errors while running. [Repry API clearly mentions that the order of exceptions mentioned in the docstring should be followed.](#) Now my implementation has been corrected according to the API.
One attack did this check and generated unexpected output from my ref monitor.

3. **Empty writeln flush:** I assumed that an empty writeln should not flush the contents of the buffer into file. I assumed this because some text editors that allow undo operation still allow undo after file has been written to disk. And the first undo operation leads to change of content and not just unsaving of the file.
But according to the assignment specifications, this assumption is wrong. The assignment spec trades off an extra undo operation for the ability to have a flush command. The [assignment spec \(second point, first bullet\)](#) mentions that the changes can be considered permanent if a valid writeln operation is performed. An empty writeln is a valid writeln according to the Repry API, thus we should allow flushing of the buffer to the file in case of empty writeln. It has been corrected now within my ref monitor.
5 attack cases did a check for this and failed due to my incorrect implementation.